B1 encoder-decoder with LoRA

CS781 Phase 2 legal named entity recognition | source-audited report

Task and evidence

Predict exact original-character entity spans in 12 legal categories. The source archive holds the actual code, configuration, run record, selected weights, and prediction ZIP. This report was produced after training from those records; it is not an executed historical notebook.

Architecture

The 6+6-layer B1 encoder-decoder supplies contextual decoder states. Additive low-rank updates train selected linear projections while the pretrained backbone weights remain frozen; a character BiGRU produces BIO tags.

Path	Legal text -> Own byte BPE -> 6-layer encoder -> 6-layer decoder + LoRA
Output	Character BiGRU -> 25 BIO logits -> Span decoder

Training and validation

Method: LoRA, rank 8. This export's archived held-out strict macro-F1 is 0.386275 on the 12-document local validation split. Its selected checkpoint and prediction ZIP hashes are in evidence/run_summary.json. The official test set has 35 document IDs; no hidden Codabench score is established by this archive.

Implementation map

Key original files: lner/models.py, lner/data.py, lner/train.py, backbones/cs781/models/b.py. The companion notebook explains the tensor route and reads archived evidence without presenting invented execution outputs.

Reproducibility boundary

The package includes exact source snapshots and small run records, plus the archived prediction ZIP. Large model weights, tokenizers, and course data stay in the separate complete archive. Original machine paths in config files need remapping for a new environment. The consensus rule was selected on validation; 100M/200M D222 exports were completed after the original deadline where marked.

Interpretation

Model variants change either pretrained backbone, parameter-update method, NER head, or pretrained checkpoint. A score difference does not isolate one factor unless all other choices match. The local held-out set is small, so its ranking should not be described as a hidden-test result.

Recorded model and adaptation

The source and run records together identify the following settings. Counts include the Phase 2 NER head. The method changes which backbone parameters are updated; it does not change the original Phase 1 tokenizer during this run.

Field	Archived value
Base family	B1
Adaptation	LoRA, rank 8
Total parameters	126553561
Trainable parameters	2272729
Backbone LR	0.0002
Head LR	0.0005
Pretrained SHA-256	d93a2f1796e90b9fdbe4…
Selected NER SHA-256	3f2b49fabeb6396549bf…

Training signal and exact-span evaluation

For the common NER runner, the character tagger minimizes weighted cross-entropy over 25 BIO categories; A0 also matches its verified B1 teacher with temperature-2 KL. D222 adds CRF negative log-likelihood and a weighted boundary objective. HeaderCRF combines cross-entropy, CRF, and boundary losses. The validation scorer counts a prediction as correct only when the start, end, and label all match the gold span; macro-F1 averages the 12 per-label F1 values equally.

Observed optimization trace

These five entries are sampled from the archived training log; loss is a per-batch training value, not held-out macro-F1. Its decline alone cannot establish improved entity extraction.

Optimizer step	Epoch	Logged batch loss
1	1	3.6160
250	2	0.6926
500	3	0.6599
750	5	0.6758
1000	6	0.5035

Label-level held-out scores

The following F1 values come from the archived selected validation prediction. Low-frequency labels are sensitive to a few documents, so read them beside their support and error counts in evidence/selection.json.

Entity label	F1	TP	FP	FN
APP	0.027	3	112	106
RESP	0.072	6	59	95
A.COUNSEL	0.348	12	12	33

Entity label	F1	TP	FP	FN
R.COUNSEL	0.388	19	48	12
JUDGE	0.333	12	25	23
WIT	0.000	0	2	21
AUTH	0.404	113	241	92
COURT	0.835	185	48	25
STAT	0.590	216	172	128
PREC	0.181	45	190	217
DATE	0.943	315	21	17
CASENO	0.513	86	111	52

Verification and reproducibility

The prediction ZIP has a single root prediction.json, passed CRC, and contains exactly 35 official test IDs. Its SHA-256 was compared with the completed run record. Each copied code file has a SHA-256 listed in evidence/SOURCE_HASHES.json. Exact re-execution requires the original selected weights, tokenizer and course data preserved in the complete archive; this compact package does not embed them or assert that historical notebook output cells existed.